

Lecture 14: MATH 342W: Introduction to Data Science and Machine Learning

Sergio E. Garcia Tapia*

March 28, 2025 (last updated March 28, 2025)

Polynomial Regression

Let $\mathcal{Y} = \mathbb{R}$ and $p = 1$. In our framework, we have

$$y = g(\mathbf{x}) + \underbrace{(h^*(\mathbf{x}) - g(\mathbf{x}))}_{\text{estimation}} + \underbrace{f(\mathbf{x}) - h^*(\mathbf{x})}_{\text{misspecification}} + \underbrace{t(\mathbf{z}) - f(\mathbf{x})}_{\text{ignorance}}$$

We have learned that we can decrease estimation error by increasing n relative to p . Reducing ignorance error requires us to gather more features. Unfortunately, that is usually not within the scope of our work. In practice, we are in control of estimation and misspecification error, but today our focus will be on the latter. The focus of this course is slowly shifting towards machine learning.

Consider the data set \mathbb{D} depicted in Figure 1 and the linear approximation h_0^* depicted, where $h_0 \in \mathcal{H}_0$ and

$$\mathcal{H}_0 = \{ w_0 + w_1 x \mid w_0, w_1 \in \mathbb{R} \}$$

We know that h_0^* is the best possible approximation to f in \mathcal{H}_0 , where f is in turn the best approximation to t with the features that we are using as proxies to the true drivers.

*Based on lectures of Dr. Adam Kapelner at Queens College. See also the [course GitHub page](#).




Figure 1: Underfitting a data set \mathbb{D} with a linear approximation h_0^* .




Figure 2: Underfitting a data set \mathbb{D} with a linear approximation h_0^* .

We reiterate that h_0^* is likely not close to f , but it is the best we can do in \mathcal{H}_0 ; simply, $f(x)$ is *not* linear in this example, so even the best linear approximation will not perform well. This is an example of **underfitting**.

Let's allow for a more *expansive* (larger) candidate set, an expanded **functional basis**. Consider:

$$\mathcal{H} = \{ w_0 + w_1x + w_2x^2 \mid w_0, w_1, w_2 \in \mathbb{R} \}$$

which includes a quadratic term, allowing us to fit parabolas and lines, and linear combinations of parabolas and lines. With this candidate set, we can better fit $f(x)$, reducing misspecification. Figure 2 shows fitting \mathbb{D} with a quadratic function $h^* \in \mathcal{H}$. Note how the coefficient w_2 of the quadratic term of the functions in \mathcal{H} are a scalar multiple of x^2 , which is the square of the single feature we have. That is, we still only measure 1 feature, but we are using it twice. In this setting, we say $p_{\text{raw}} = 1$, and $p = 2$; we call x^2 a **transformed** or **derived** feature. Compare this with using a linear model, where we would introduce a new coefficient b_2 , but in that case, it would be an example of chance capitalization, and hence overfitting.

Using derived features that are squares, cubes, etc, of the raw features is an example of **polynomial regression**. Here's a question: is

$$\hat{y} = b_0 + b_1x + b_2x^2$$

a linear model? We can justify the answer both ways:

- **Yes:** It is linear because \mathcal{H} is a set of linear combinations.
- **No:** It is not linear because it is not a straight line.

Now, we may ask whether it is justified or “principled” to use polynomial regression. In fact, it is, and it is a consequence of the famous Stone-Weierstrass theorem, which proves that any continuous function $f : \mathbb{R}^p \rightarrow \mathbb{R}$ can be approximated by a sum of polynomials. If you have learned basic calculus, you may have seen a similar example of this: Taylor's Theorem. Indeed, Taylor Theorem says that an infinitely differentiable function can be expressed as an infinite sum of polynomials. However, Stone-Weierstrass does not require differentiability, only continuity.

How do we find $g \in \mathcal{H}$ in the case of polynomial regression? Suppose we have $\mathbb{D} = \{(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)\}$. Then we can transform our matrix of raw features into a design matrix X that has the derived features used in quadratic regression:

$$X_{\text{raw}} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \implies X = \begin{bmatrix} 1 & x_1 & x_1^2 \\ 1 & x_2 & x_2^2 \\ \vdots & \vdots & \vdots \\ 1 & x_n & x_n^2 \end{bmatrix}$$

In the R programming language, which would simply be:

```
X = cbind(1, x, x^2)
```

Then we can *apply OLS using matrix X, just like before*:

$$\begin{aligned} \mathbf{b} &:= (X^\top X)^{-1} X^\top \mathbf{y} \\ \mathbf{e} &:= \mathbf{y} - \hat{\mathbf{y}} = \mathbf{y} - X\mathbf{b} \\ SSE &:= \|\mathbf{e}\|^2 \\ MSE &:= \frac{SSE}{n - (p + 1)} \\ RMSE &:= \sqrt{MSE} \end{aligned}$$

Note that we are making a tradeoff. By allowing a quadratic term, our misspecification error increases, but our estimation error increases also. If $p+1 \lll n$, then increasing p via these derived features may be justified if the misspecification error decreases meaningfully.

The Design Matrix with Transformed Features

One concern we need to address is whether the matrix X obtained from X_{raw} is full rank after adding the new column whose entries are squares of the raw feature. One way to justify it is as follows. Suppose that the columns are linearly dependent. Then there are coefficients c_0, c_1, c_2 making up a vector $\mathbf{c} = [c_0 \ c_1 \ c_2]^\top \in \mathbb{R}^3$ such that $X\mathbf{c} = \mathbf{0}_{n \times 1}$. If we define a polynomial $p(x) = c_0 + c_1x + c_2x^2$, then

$$\begin{bmatrix} 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix} = \mathbf{0}_{n \times 1} = X\mathbf{c} = \begin{bmatrix} 1 & x_1 & x_1^2 \\ 1 & x_2 & x_2^2 \\ \vdots & \vdots & \vdots \\ 1 & x_n & x_n^2 \end{bmatrix} \begin{bmatrix} c_0 \\ c_1 \\ c_2 \end{bmatrix} = \begin{bmatrix} c_0 + c_1x_1 + c_2x_1^2 \\ c_0 + c_1x_2 + c_2x_2^2 \\ \vdots \\ c_0 + c_1x_n + c_2x_n^2 \end{bmatrix} = \begin{bmatrix} p(x_1) \\ p(x_2) \\ \vdots \\ p(x_n) \end{bmatrix}$$

Hence, $p(x_1) = p(x_2) = \dots = p(x_n) = 0$. However, since p is a polynomial of degree two, it has at most 2 zeroes. Therefore, if at least 3 distinct values for the feature x exist in \mathbb{D} , it's impossible for p to be zero on all of them. In general we deal with $p + 1 \lll n$, so that we can usually guarantee X is full rank.

In particular, if X has binary features (that is, its values can be only 0 and 1), then we cannot apply this transformation, since squaring will result in the same value, hence yielding a linearly dependent column.

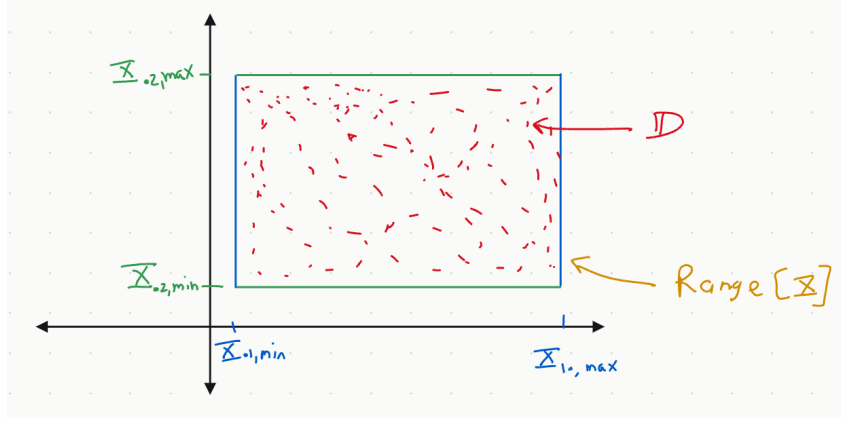


Figure 3: A depiction of the range of X , as described in Definition .

Interpretability and Machine Learning

One point is whether our new result is interpretable, and how much that matters to us. Consider

$$\hat{y} = b_0 + b_1x + b_2x^2 = b_0 + (b_1 + b_2x)x$$

We can say that if x changes by 1, then \hat{y} changes by $b_1 + b_2x$. What if we go further, fitting a cubic?

$$\hat{y} = g(x) = b_0 + b_1x + b_2x^2 + b_3x^3$$

As we fit higher degree polynomials, we lose easy the interpretation that we get from a linear model (and quadratic to a lesser extent). However, we do get a closer fit. The tradeoff is interpretability and model complexity.

This matters in machine learning, which we now define.

Definition (Machine Learning). Machine learning is supervised learning (some people stop here), with a very flexible \mathcal{H} .

A downside of a very flexible \mathcal{H} is the loss of interpretability.

Interpolation and Extrapolation

Definition. Given design matrix X , we define the **range of X** to be a rectangle in \mathbb{R}^p :

$$\text{Range}[X] = [X_{\cdot 1, \min}, X_{\cdot 1, \max}] \times [X_{\cdot 2, \min}, X_{\cdot 2, \max}] \times \cdots \times [X_{\cdot p, \min}, X_{\cdot p, \max}]$$

Here, \times is used to denote a Cartesian product, $X_{\cdot j}$ denotes the j th column of X (containing all n instances of the j th feature), and $[X_{\cdot j, \min}, X_{\cdot j, \max}]$ denotes a closed interval in \mathbb{R} , delimited by the smallest value of the j th feature and the largest value of the j feature in X (see Figure 3).

Definition. Given a design matrix X :

- **Interpolation** is predicting for $\mathbf{x} \in \text{Range}[X]$.
- **Extrapolation** is predicting for $\mathbf{x} \notin \text{Range}[X]$.

An important fact to remember is that models are defined for *interpolation*, and bad things happen during extrapolation. Interestingly, different candidate sets \mathcal{H} and different algorithms \mathcal{A} lead to different behavior under extrapolation.

Logarithmic Approximations

At times, it can be useful to perform a transformation on the response and (or) the feature values. A popular such transformation is a *logarithmic transformation*. First, recall the following Taylor series approximation:

$$\ln(1+x) = x - \frac{x^2}{2} + \frac{x^3}{3} - \cdots + \cdots \approx x, \quad \text{if } x \approx 0$$

Therefore,

$$\ln((1+x) - 1) \approx x - 1 \tag{1}$$

We will apply (1) in the following discussion.

\hat{y} vs $\log(x)$

Now suppose we transform a feature x into $\ln(x)$, and apply OLS, yielding the model

$$\hat{y} = b_0 + b_1 \ln(x)$$

We want to try to interpret what happens if x changes. Let $\Delta x := x_f - x_0$, where x_0 is the initial value of x and x_f is the final value of x after some change. Then the change in the prediction \hat{y} is:

$$\begin{aligned} \Delta \hat{y} &= (b_0 + b_1 \ln(x_f)) - (b_0 + b_1 \ln(x_0)) \\ &= b_1 (\ln(x_f) - \ln(x_0)) \\ &= b_1 \ln \left(\frac{x_f}{x_0} \right) \\ &\approx b_1 \left(\frac{x_f}{x_0} - 1 \right) && \text{(by (1))} \\ &= b_1 \underbrace{\left(\frac{x_f - x_0}{x_0} \right)}_{\text{proportional change}} \end{aligned}$$

We can interpret this as follows: if x increases by 27%, then the predicted response increases by $b_1 \cdot 27\%$.

$\log(\hat{y})$ vs x

Next, suppose we transform the response into $\ln(y)$ and leave the feature x unchanged:

$$\ln(\hat{y}) = b_0 + b_1 x$$

Then if x changes by $\Delta x = x_f - x_0$, the corresponding change in the $\ln(\hat{y})$ is

$$\begin{aligned}\Delta \ln(\hat{y}) &= \ln(\hat{y}_f) - \ln(\hat{y}_0) \\ &= (b_0 + b_1 x_f) - (b_0 + b_1 x_0) \\ &= b_1 \Delta x\end{aligned}$$

Hence,

$$\begin{aligned}b_1 \Delta x &= \ln\left(\frac{\hat{y}_f}{\hat{y}_0}\right) \\ &\approx \frac{\hat{y}_f - \hat{y}_0}{\hat{y}_0} \quad (\text{by (1)})\end{aligned}$$

The interpretation is that if there is a one unit change in x , then the predicted response is a b_1 proportion change. For example, if $b_1 = 0.37$, then \hat{y} changes by 37%.

$\log(y)$ vs $\log(x)$

Suppose we transform both the feature and the response by applying the logarithm:

$$\ln(\hat{y}) = b_0 + b_1 \ln(x)$$

Then a similar calculation yields

$$\frac{\hat{y}_f - \hat{y}_0}{\hat{y}_0} \approx b_1 \frac{x_f - x_0}{x_0}$$

Let $b_1 \approx 0.37$. If x increases by 27%, then the predicted response increases by $0.37 \cdot 0.27 \cdot 100$ percent.

Summary on Transformations

Usually, transformations are done on the features (the x 's). We saw one example in which we transform the response y to $\ln(y)$. In fact, the most common transformations for $\mathcal{Y} = \mathbb{R}$ are \ln , \log_{10} , and \log_2 . Note that doing this amounts to the following:

$$\mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} \implies \mathbf{y}' = \begin{bmatrix} \ln(y_1) \\ \ln(y_2) \\ \vdots \\ \ln(y_n) \end{bmatrix}$$

and then using OLS as usual with $\mathbf{y} = (X^\top X)^{-1} X^\top \mathbf{y}'$. This yields

$$\begin{aligned}\ln(\hat{y}) &= b_0 + b_1 x_1 + \cdots + b_p x_p \iff \\ \hat{y} &= e^{b_0} e^{b_1 x_1} \cdots e^{b_p x_p} \\ &= e^{b_0} (e^{b_1})^{x_1} \cdots (e^{b_p})^{x_p} \\ &= m_0 \cdot m_1^{x_1} \cdots m_p^{x_p}\end{aligned}$$




Figure 4: A histogram of the responses values y .




Figure 5: A logarithmic transformation applied to the response y from Figure 4

which becomes a multiplicative model. We can do this, but we must bear in mind that the units of the response change. Thus, we eventually must reverse the transformation. Indeed, we cannot compare the response of the original model with that of the transformed model, because the units do not match. We could use R^2 since it is unitless, but since the R^2 metric is hard to interpret, it is often not as useful.

Another thing to bear in mind is that if we apply a logarithmic transformation, then a small residual in the logarithmic change can still indicate a large error in the original units, because we must apply the exponential function to reverse it.

In summary, we can transform the response, as long as we are careful about the units of the predictions, the error metrics, and how we compare model fits.

One good reason to perform a transformation is that y may be right-skewed, which can make it difficult to capture extremes, as depicted in Figure 4.

Applying a logarithmic transformation results in something like in Figure 5.

Generalized Additive Model (GAM)

Imagine adding many transformations to your design matrix X . For example, polynomial, logarithmic, trigonometric functions, etc. Your model may look like

$$g(\mathbf{x}) = g_1(x_1) + g_2(x_2) + \cdots + g_p(x_p)$$

where the g_j 's are arbitrarily complex and possibly non-linear (note that in simple polynomial regression, we have $g_j(x) = x^j$). This is called a **generalized linear model**

(**GAM**). Although well-studied, GAMs have fallen out of favor, as we will see when we discuss random forests.

What is missing from this type of model \mathcal{H} ? Functions of multiple variables. Notice that functions of multiple variables are missing from the type of model \mathcal{H} implied by GAM, so we have

$$\frac{\partial}{\partial x_\ell}[g_j(x_j)] = 0, \quad \forall_{j \neq \ell}$$

Put another way, the model does not account for *interactions between features*. The following model for $p = 2$ allows for a *first-order interaction*:

$$\mathcal{H} = \{ w_0 + w_1x + w_2x_2 + w_3x_1x_2 \mid \mathbf{w} \in \mathbb{R}^4 \}$$

Here, x_1 interacts with x_2 , and the transformation from X_{raw} is

$$X_{\text{raw}} = \begin{bmatrix} x_{11} & x_{21} \\ x_{12} & x_{22} \\ \vdots & \vdots \\ x_{1n} & x_{2n} \end{bmatrix} \implies X = \begin{bmatrix} 1 & x_{11} & x_{21} & x_{11}x_{21} \\ 1 & x_{12} & x_{22} & x_{12}x_{22} \\ \vdots & \vdots & \vdots & \vdots \\ 1 & x_{1n} & x_{2n} & x_{1n}x_{2n} \end{bmatrix}$$

Again, we apply OLS to get $\mathbf{b} = (X^\top X)^{-1}X^\top \mathbf{y}$. This time, the prediction is given by

$$\begin{aligned} \hat{\mathbf{y}} &= b_0 + b_1x_1 + b_2x_2 + b_3x_1x_2 \\ &= b_0 + b_1x_1 + (b_2 + b_3x_1)x_2 \\ &= b_0 + b_2x_2 + (b_1 + b_3x_2)x_1 \end{aligned}$$

We can interpret it the following way:

- If x_2 increases by 1 unit, then \hat{y} changes by $b_2 + b_3x_1$.
- If x_1 increases by 1 unit, then \hat{y} changes by $b_1 + b_3x_2$.