

Lecture 21: MATH 342W: Introduction to Data Science and Machine Learning

Sergio E. Garcia Tapia*

April 29th, 2025 (last updated May 13, 2025)

Bagging

Recall bagging, depicting in Figure 1. How do we do validation for bagged models? Consider the first bootstrap sample \mathbb{D}_1 , and the model we compute on it:

$$g_1 = \mathcal{A}(\mathbb{D}_1, \mathcal{H})$$

Then $\mathbb{D} \setminus \mathbb{D}_1$ is out-of-sample for g_1 , meaning the metrics computed on $\mathbb{D} \setminus \mathbb{D}_1$ are honest. We can do the same for g_2 computed on bootstrap sample \mathbb{D}_2 (likely different 1/3 of data missing than \mathbb{D}_1). We call this metric **out-of-bag (oob)**. If we do this M times, we have M bootstrap samples $\mathbb{D}_1, \dots, \mathbb{D}_M$, each missing a slightly different 1/3 of the data (see Figure 2), and we compute corresponding prediction functions g_1, \dots, g_M . For each unit in \mathbb{D} , we can collect the models for which that unit is out-of-bag and use it to compute an out-of-bag metric for g_{avg} which is given by

$$g_{avg} := \frac{1}{M} \sum_{m=1}^M g_m$$

*Based on lectures of Dr. Adam Kapelner at Queens College. See also the [course GitHub page](#).




Figure 1: Depiction of bagging. Each bootstrapped sample contains about $\frac{2}{3}n$ of the original data.




Figure 2: Out-of-bag data in bootstrap samples. Each circle depicts which data point is missing from a given bootstrap sample.

Example. Given M bootstrap models with index set $\{1, 2, \dots, M\}$, suppose models 17, 37, 88 do not contain the observation $i = 3$ (i.e., the data point (\mathbf{x}_3, y_3) is out of bag for them). Then we can use models g_{17}, g_{37}, g_{88} to compute $\hat{y}_{17}, \hat{y}_{37}, \hat{y}_{88}$ on \mathbf{x}_3 . Thus, the out-of-bag value of g_{avg} for observation $i = 3$ is

$$\hat{y}_{oob,3} = \text{Average}(\hat{y}_{17}, \hat{y}_{37}, \hat{y}_{88})$$

We do this computation for all $i = 1, 2, \dots, n$. In the sample above, the out-of-bag prediction was computed using 3 values. However, you will generally observe about $\frac{1}{3}M$ values included in the out-of-bag calculation.

Bagging is usually done with trees ($N_0 = 1$) because bias is low.

Random Forests

Recall ρ is the correlation between the trees. Thus, reducing ρ means reducing coupling (correlation) between the trees. In 2001, Breiman brought the ρ term in the Bias-Variance decomposition formula down even further. Breiman suggested that, instead of doing a full greedy search (as is done in regression trees), we can do the following:

Algorithm 1 (Random Forest Algorithm). Use the bagged tree model with one change to the CART algorithm. Pick $m_{\text{try}} < p_{\text{raw}}$ and only search a random subset of $\{1, 2, \dots, p_{\text{raw}}\}$ of size m_{try} at each node. The “reasonable” defaults for m_{try} are $m_{\text{try}} = \lfloor p_{\text{raw}}/3 \rfloor$ for regression, and $m_{\text{try}} = \lfloor \sqrt{p_{\text{raw}}} \rfloor$ for classification.

Notice that the transition to random forests involves a change in the algorithm (\mathcal{A}), not \mathcal{H} or the features.

Example. Let $p_{\text{raw}} = 10$, $m_{\text{try}} = 3$. For each root node, sample a subset of $m_{\text{try}} = 3$. For example, $\{x_3, x_7, x_9\}$, and do a greedy split based on those candidate splits only.

If m_{try} is p_{raw} , then we have done nothing, since we would have the original regression

tree algorithm. If m_{try} is too low (say, $m_{\text{try}} = 1$), then we risk having a bias that is too large due to underfitting. In spite of the “defaults” in place, we really should use model selection to optimize m_{try} , which is effectively a hyperparameter.

In summary, random forests contain two types of randomness: in the units (from bagging) and in the features (from m_{try}). The term forest comes from the fact that we have multiple trees.

R Demos

See `QC_MATH_342W_Spring_2025/practice_lectures/lec21.Rmd`.

Missingness

Suppose we have a data set \mathbb{D} for which some units have missing features. That is, if we make a matrix X containing the values of the features of all the units, then there are missing entries.

In the R programming language, these would be `NA` values. Note that there is no missingness in the response y (otherwise, it would not be in \mathbb{D}). If we have missing data, we cannot do this:

$$\mathbf{b} = (X^\top X)^{-1} X^\top \mathbf{y}$$

For any of the algorithms \mathcal{A} that we have discussed, we cannot have missing data. The simplest strategy is called **listwise deletion**, where we drop all subjects (rows) from \mathbb{D} where there is at least one missing value. You should only do this if n is large and the number of missingness rows is “very small”. A down side of this is it incurs more estimation. In the following lecture, we will discuss another problem with listwise deletion, and discuss other strategies for handling missing data.